



Distributed Deep Learning

Huihuo Zheng

Data Science Team
Argonne Leadership Computing Facility

huihuo.zheng@anl.gov

www.anl.gov

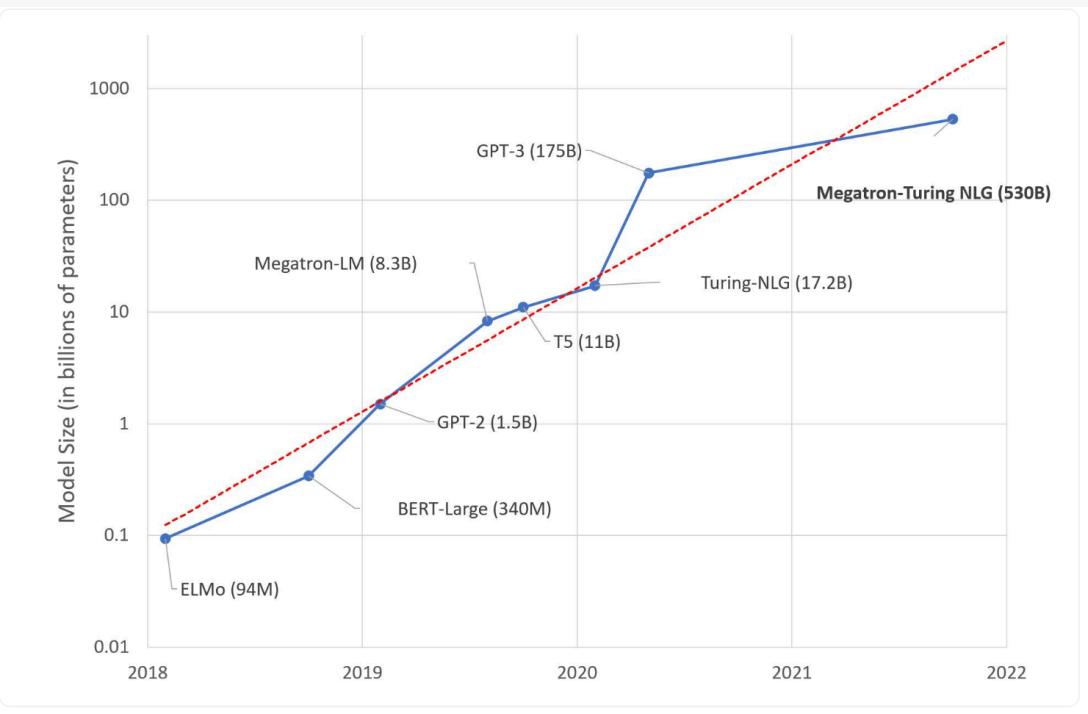
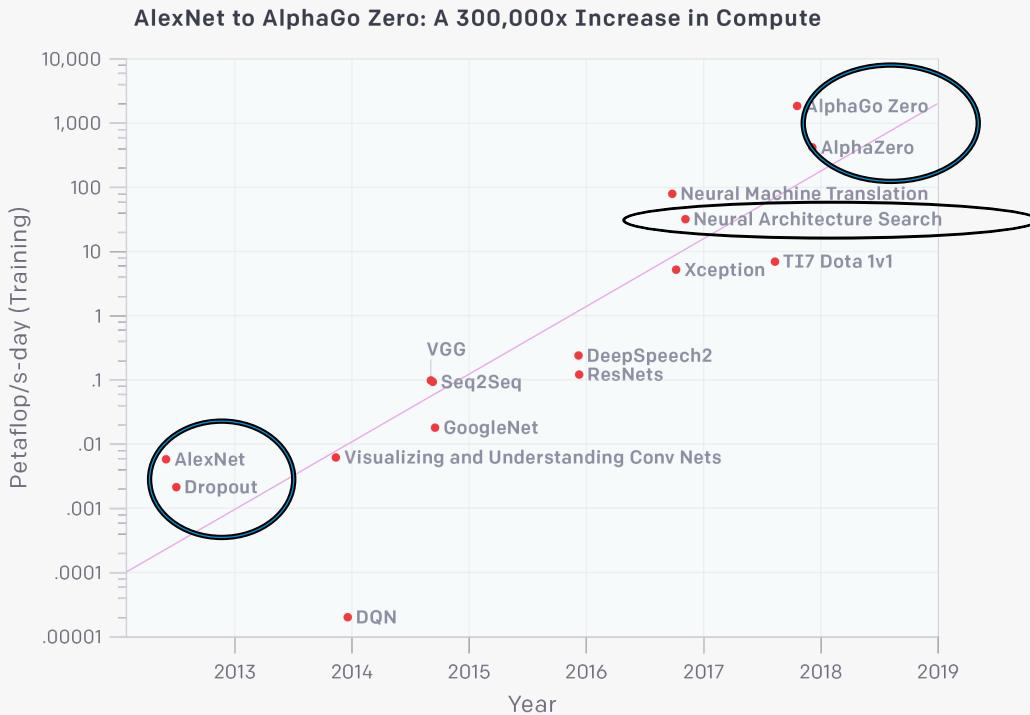
Presented by Varuni Sastry
vsastry@anl.gov

Outline

- The need for distributed training
- State-of-the-art parallelization schemes
- Data parallel training in details
- I/O and data management in distributed training
- Hands on

The need for distributed training on HPC

*“Since 2012, the amount of compute used in the largest AI training runs has been increasing exponentially with a **3.5 month** doubling time (by comparison, Moore’s Law had an **18 month** doubling period).”*



<https://openai.com/blog/ai-and-compute/>

Large language model: # parameters grows by about 10x every year

Distributed deep learning for ResNet-50

TRAINING TIME AND TOP-1 VALIDATION ACCURACY WITH RESNET-50 ON IMAGENET

		Batch Size	Processor	DL Library	Time	Accuracy
He et al. [1]	2016	256	Tesla P100 × 8	Caffe	29 hours	75.3 %
Goyal et al. [2]		8,192	Tesla P100 × 256	Caffe2	1 hour	76.3 %
Smith et al. [3]		8,192 → 16,384	full TPU Pod	TensorFlow	30 mins	76.1 %
Akiba et al. [4]		32,768	Tesla P100 × 1,024	Chainer	15 mins	74.9 %
Jia et al. [5]		65,536	Tesla P40 × 2,048	TensorFlow	6.6 mins	75.8 %
Ying et al. [6]		65,536	TPU v3 × 1,024	TensorFlow	1.8 mins	75.2 %
Mikami et al. [7]		55,296	Tesla V100 × 3,456	NNL	2.0 mins	75.29 %
Yamazaki et al	2019	81,920	Tesla V100 × 2,048	MXNet	1.2 mins	75.08%

Quoted from Masafumi Yamazaki, arXiv:1903.12650

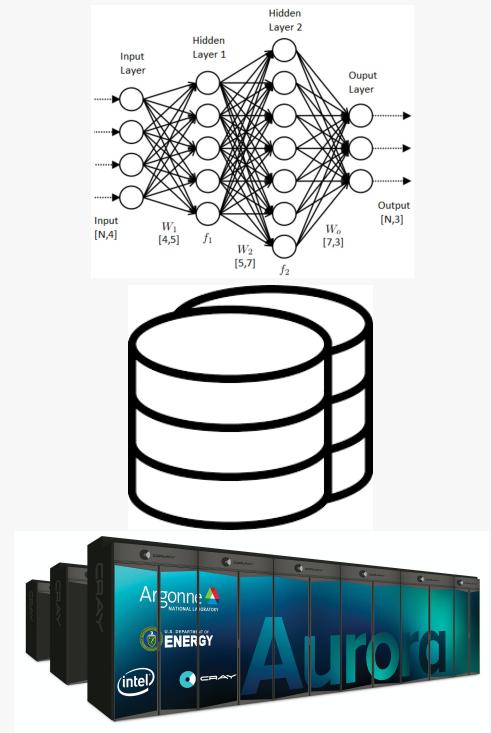
Training Large Natural Language Model is expensive

Scheme	Number of parameters (billion)	Model-parallel size	Batch size	Number of GPUs	Microbatch size	Achieved teraFLOP/s per GPU	Training time for 300B tokens (days)
ZeRO-3 without Model Parallelism	174.6	1	1536	384	4	144	90
				768	2	88	74
				1536	1	44	74
	529.6	1	2560*	640	4	138	169
				1120	2	98	137
				2240	1	48	140
PTD Parallelism	174.6	96	1536	384	1	153	84
				768	1	149	43
				1536	1	141	23
	529.6	280	2240	560	1	171	156
				1120	1	167	80
				2240	1	159	42

Narayanan, D et al. Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*; ACM: St. Louis Missouri, 2021; pp 1–15.

The need for distributed training

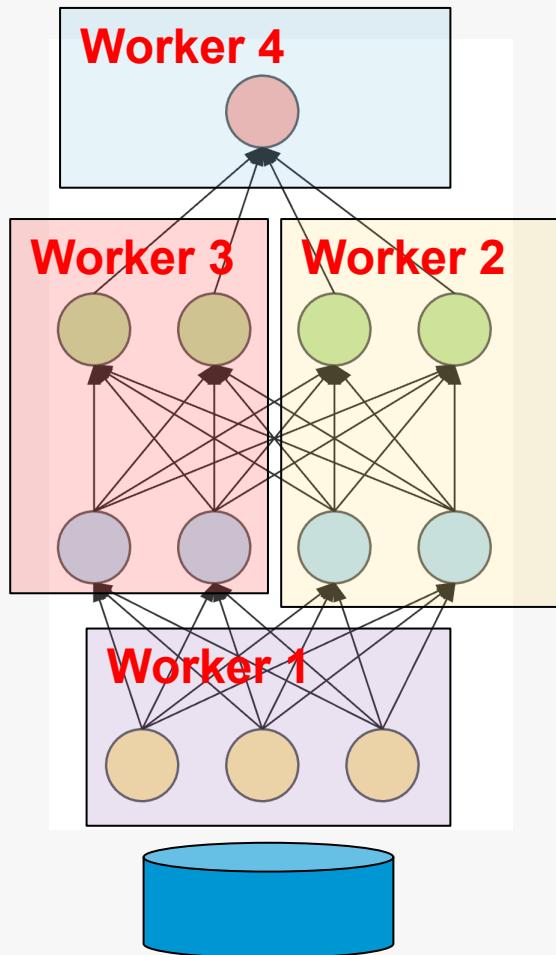
- Increase of model complexity leads to dramatic increase of the amount of computation;
- Increase of the size of dataset makes sequentially scanning the whole dataset increasingly impossible;
- Coupling of deep learning to traditional HPC simulations might require distributed training and inference.



Examples of scientific large scale deep learning

- Thorsten Kurth, Exascale Deep Learning for Climate Analytics, arXiv:1810.01993 (Gordon Bell Prize)
- R. M. Patton, Exascale Deep Learning to Accelerate Cancer Research, arXiv:1909.1229
- N. Laanait, Exascale Deep Learning for Scientific Inverse Problems, arXiv:1909.11150
- W. Dong et al, Scaling Distributed Training of Flood-Filling Networks on HPC Infrastructure for Brain Mapping, arXiv:1905.06236
- A Khan, et al, Deep learning at scale for the construction of galaxy catalogs in the Dark Energy Survey Physics Letters B 795, 248-258
- Narayanan, D.; et al, Efficient Large-Scale Language Model Training on GPU Clusters Using Megatron-LM. arXiv: [2104.04473](https://arxiv.org/abs/2104.04473)

Parallelization schemes – Model Parallelism (MP)



Model parallelism

```
import torch
import torch.nn as nn
import torch.optim as optim

class ToyModel(nn.Module):
    def __init__(self):
        super(ToyModel, self).__init__()
        self.net1 = torch.nn.Linear(10, 10).to('cuda:0')
        self.relu = torch.nn.ReLU()
        self.net2 = torch.nn.Linear(10, 5).to('cuda:1')

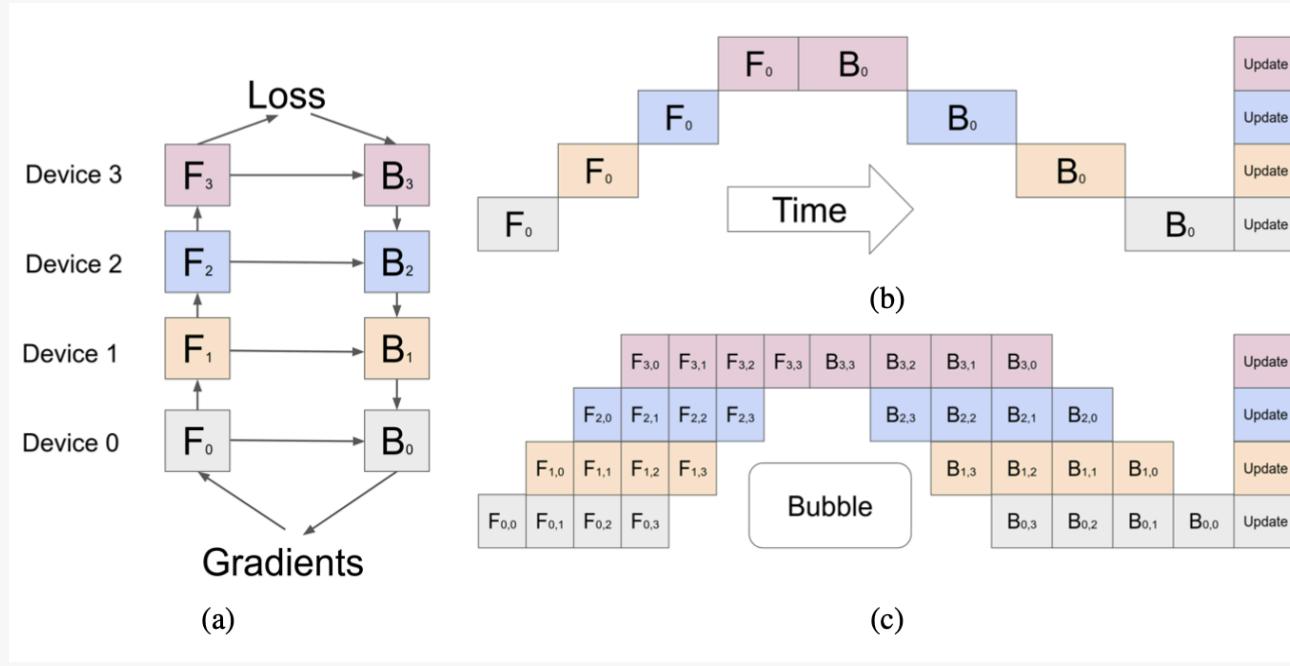
    def forward(self, x):
        x = self.relu(self.net1(x.to('cuda:0')))
        return self.net2(x.to('cuda:1'))

model = ToyModel()
loss_fn = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.001)

optimizer.zero_grad()
outputs = model(torch.randn(20, 10))
labels = torch.randn(20, 5).to('cuda:1')
loss_fn(outputs, labels).backward()
optimizer.step()
```

PyTorch multiple GPU
model parallelism
within a node

Parallelization schemes – Pipeline parallelism (PP)

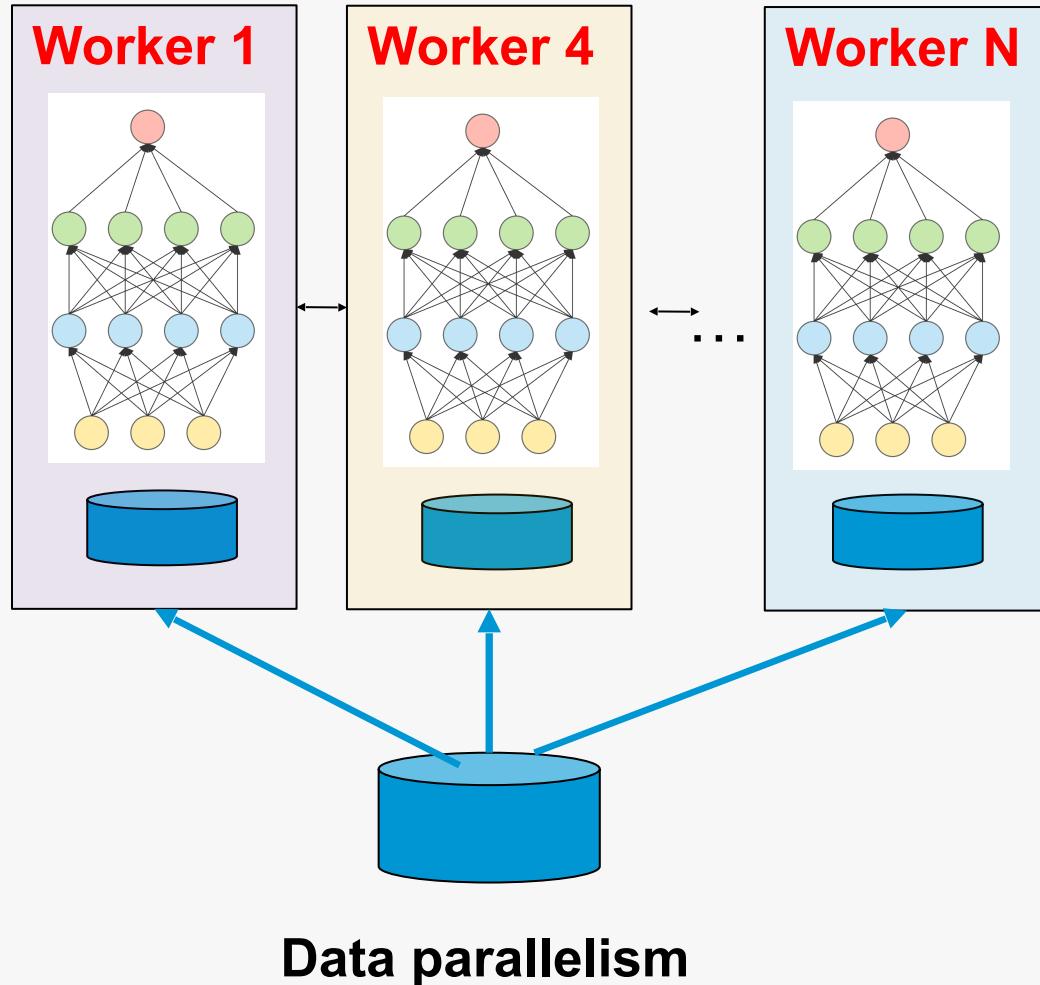


Pipeline libraries:

- GPipe: arXiv:1811.06965
- Pipe-torch: DOI: 10.1109/CBD.2019.00020
- PipeDream: arXiv:1806.03377
- HetPipe: arXiv:2005.14038
- DAPPLE: arXiv:2007.01045
- PyTorch Distributed RPC Frameworks:
https://pytorch.org/tutorials/intermediate/dist_pipeline_parallel_tutorial.html
- DeepSpeed: <https://github.com/microsoft/DeepSpeed>

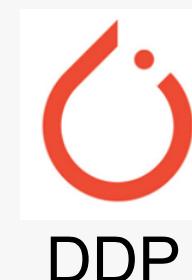
- Partition model layers into multiple groups (stages) and place them on a set of inter-connected devices.
- Each input batch is further divided into multiple micro-batches, which are scheduled to run over multiple devices in a pipelined manner.

Parallelization schemes – Data Parallelism (DP)



- Model is replicated on each worker
- Each worker processes a subset of the minibatch
- Sync up the weights before updating the model

Popular frameworks support DP



<https://leimao.github.io/blog/PyTorch-Distributed-Training/>
<https://github.com/horovod/horovod>
<https://github.com/microsoft/DeepSpeed>

Data parallel training in details

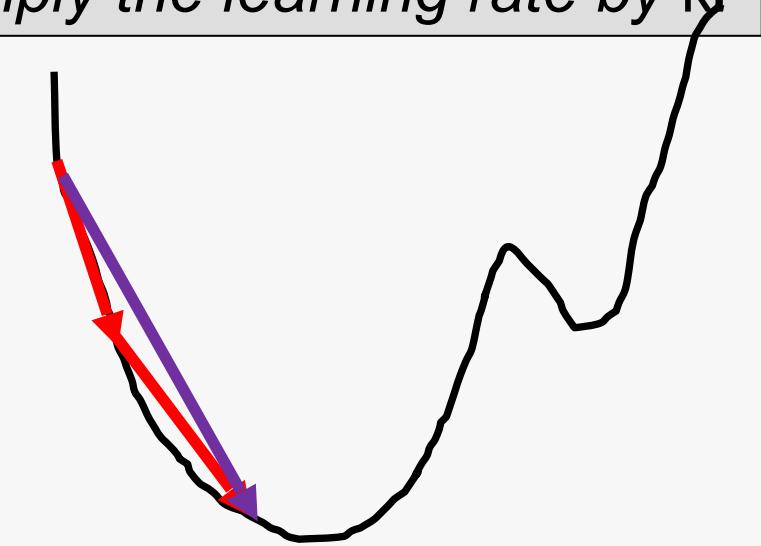
Linear rate scaling rule

When the minibatch size is multiplied by k, multiply the learning rate by k.

Mini-batch stochastic Gradient Descent

$$w_{t+1} = w_t - \eta \frac{1}{n} \sum_{x \in \mathcal{B}} \nabla l(x, w_t)$$

Learning rate, lr Mini-batch



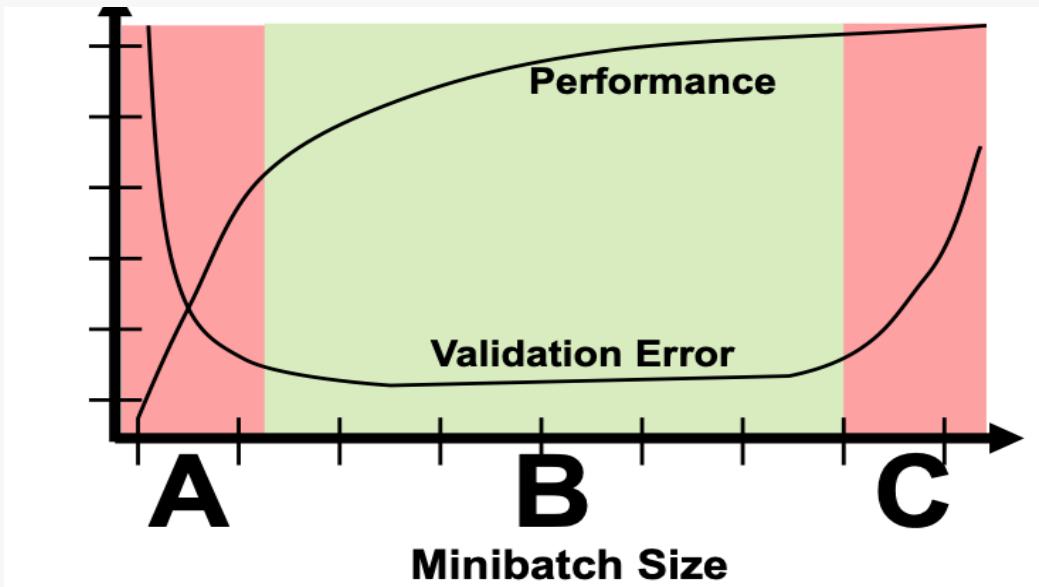
$$\text{lr}_{\text{scale}} = \text{lr}^* \text{nprocs}$$

Typical practice / suggestion:

- keep local batch size per worker, i.e., increase the global batch size linearly
- Increase the learning rate proportionally

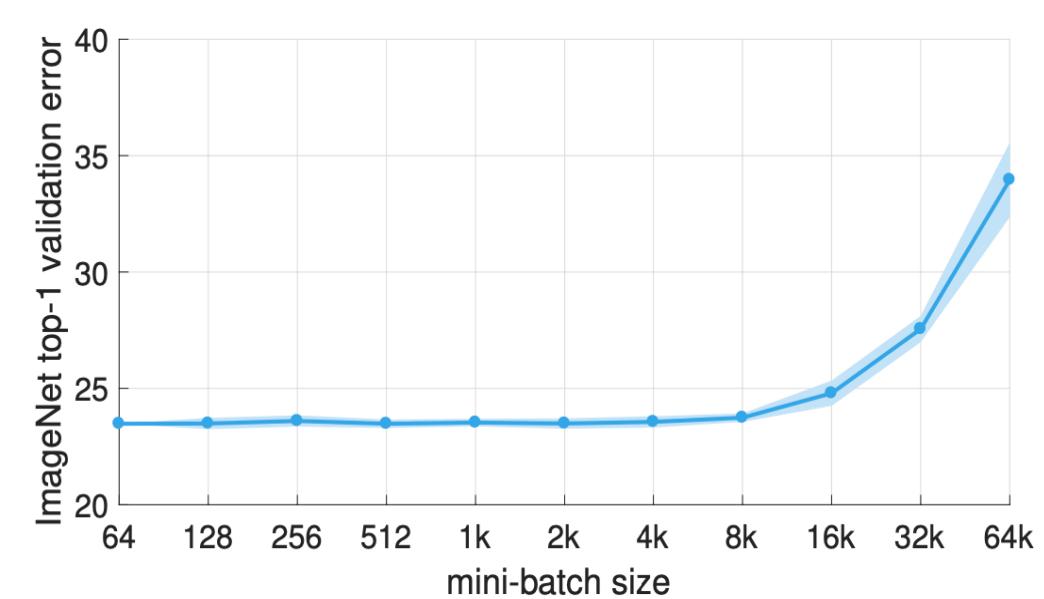
Actuality: need to adjust the batch size and learning rate at different scale

Potential issues for large batch size and learning rate



Tal Ben-Nun and Torsten Hoefler, arXiv:1802.09941

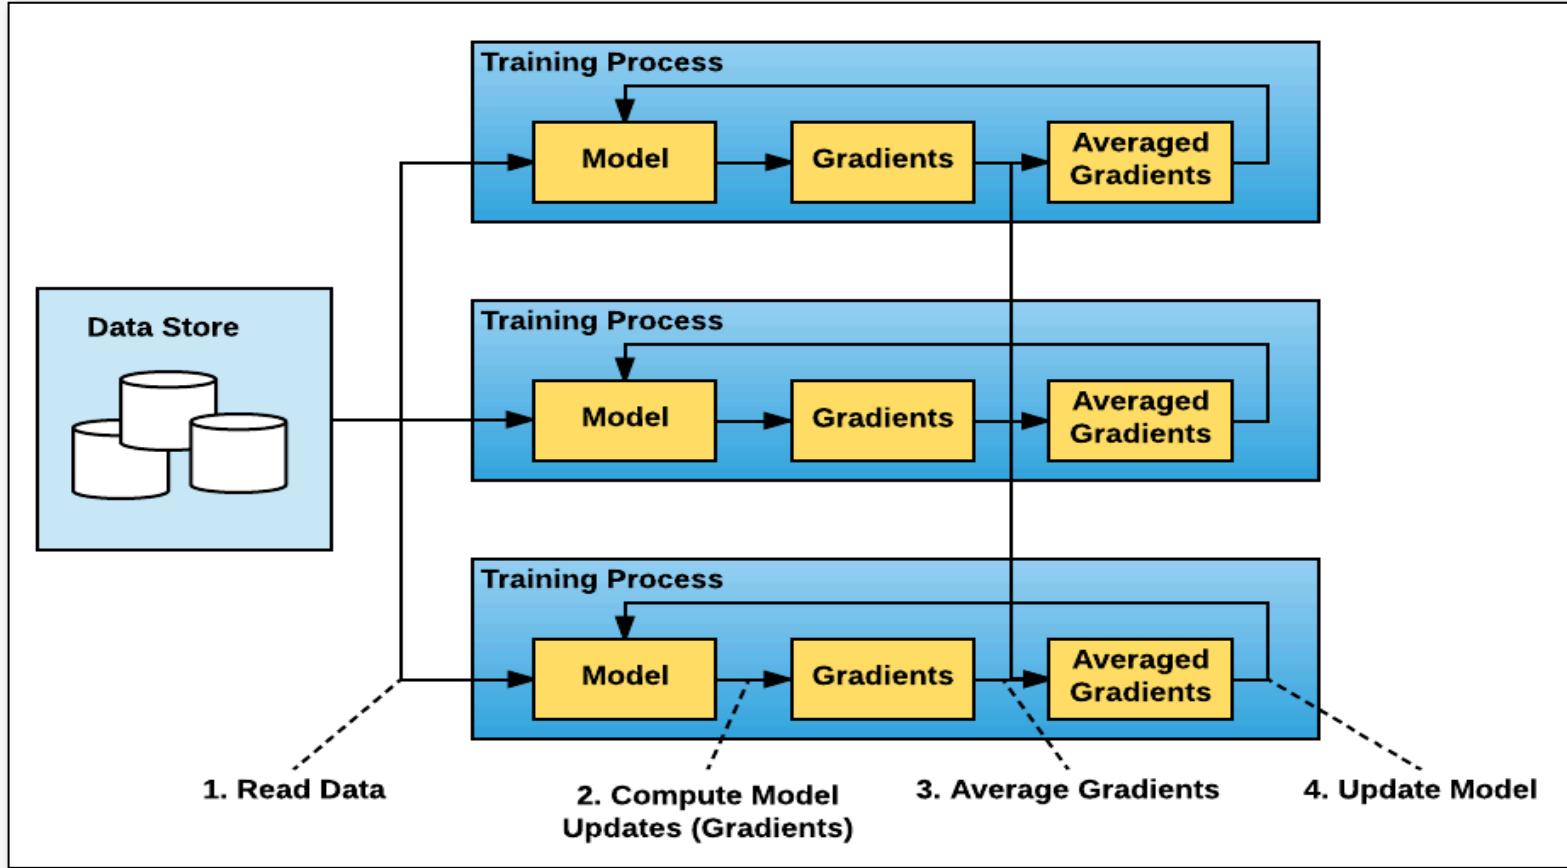
- Stability of optimization – can be solved by learning rate warming up
- Generation gap – systematic issue; has to reduce the learning rate



Validation error for different mini-batch size for Resnet50

P. Goyal et al, arXiv: 1706.02677

Data parallel training with Horovod



<https://eng.uber.com/horovod/>

Data parallel training with Horovod

Steps to parallelize a series code:

- Import Horovod modules and initialize horovod
- Scale the learning rate by number of workers
- Wrap optimizer in hvd.DistributedOptimizer
- Broadcast the weights from worker 0 to all the workers
- Worker 0 saves the check point files
- Dataset sharding: make sure different workers load different samples.

Instruction on how to change the code is [here](#)

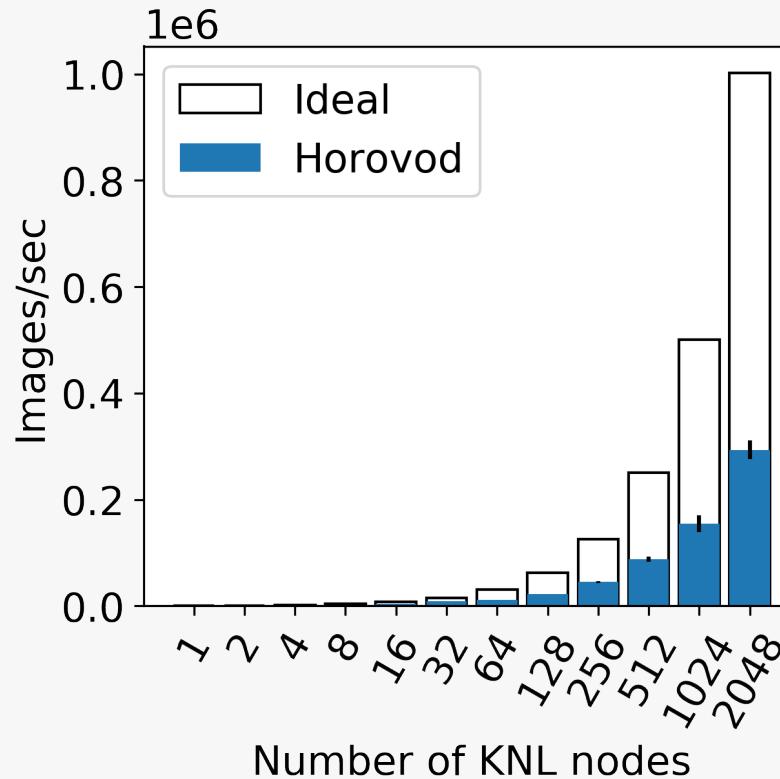
Tensorflow: [04_keras_cnn_verbose_hvd.py](#)

Pytorch: [04_pytorch_cnn_hvd.py](#)

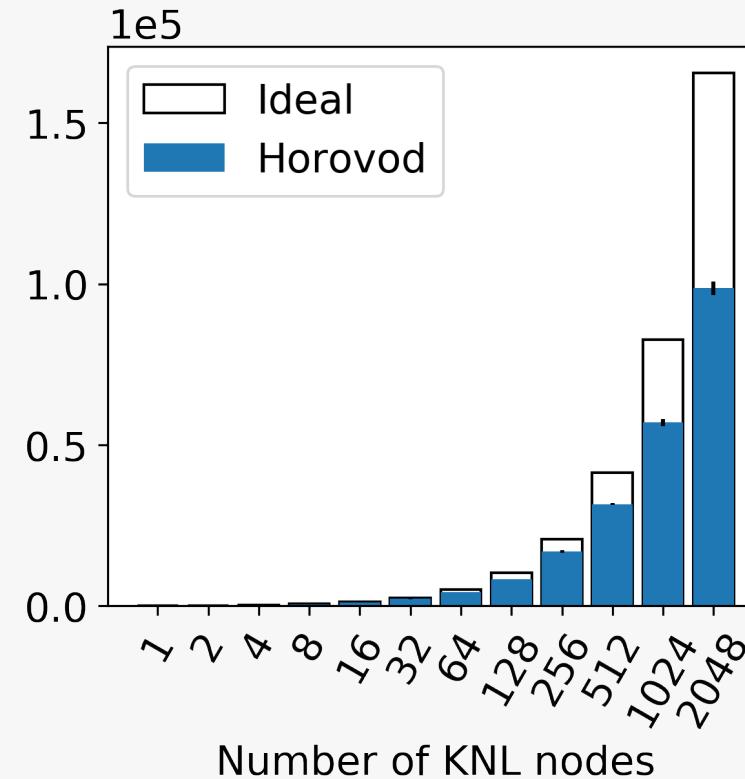


<https://eng.uber.com/horovod/>

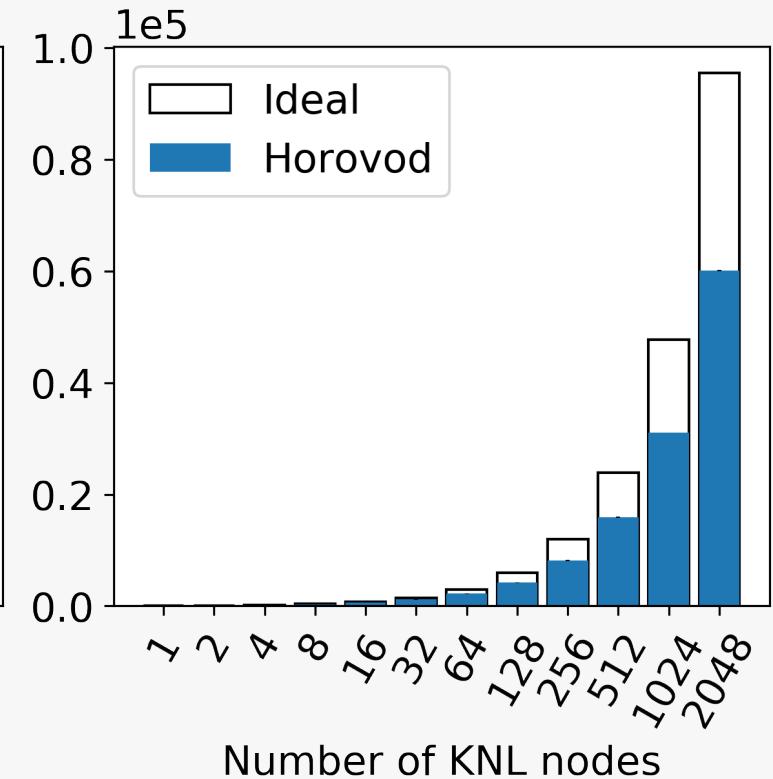
Scaling TensorFlow using Data parallelism on Theta @ ALCF: fixing local batch size = 512



AlexNet

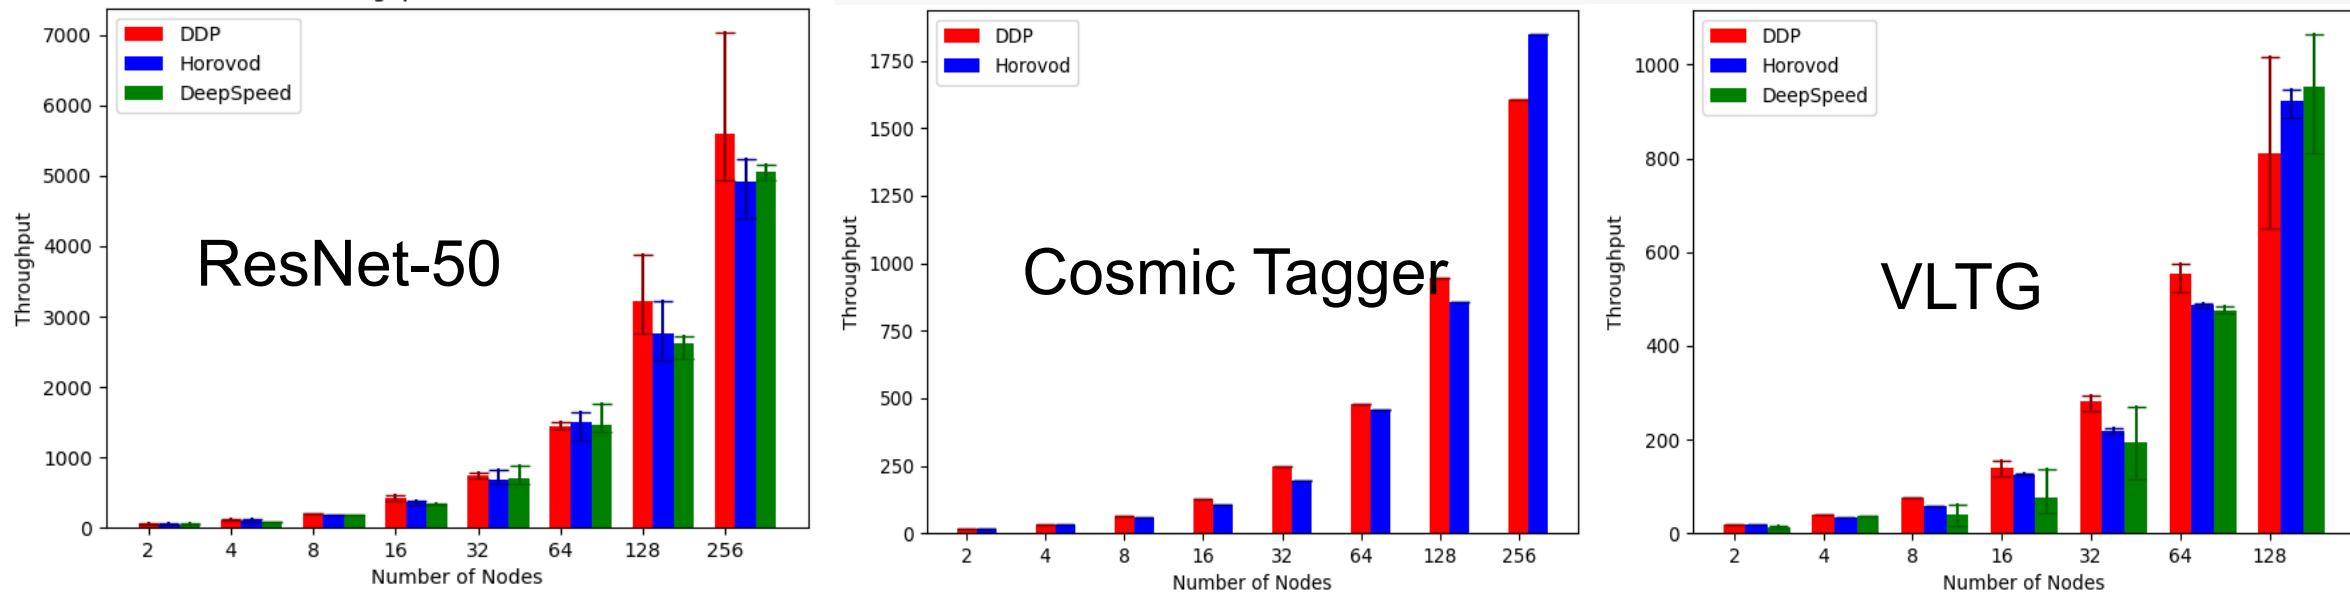


ResNet-50



Inception V3

Different frameworks show similar scaling efficiency



Horovod, DDP, and DeepSpeed show similar performance for three PyTorch models. Evaluation were done on Polaris @ ALCF



Zhenhao Z.
@IIT

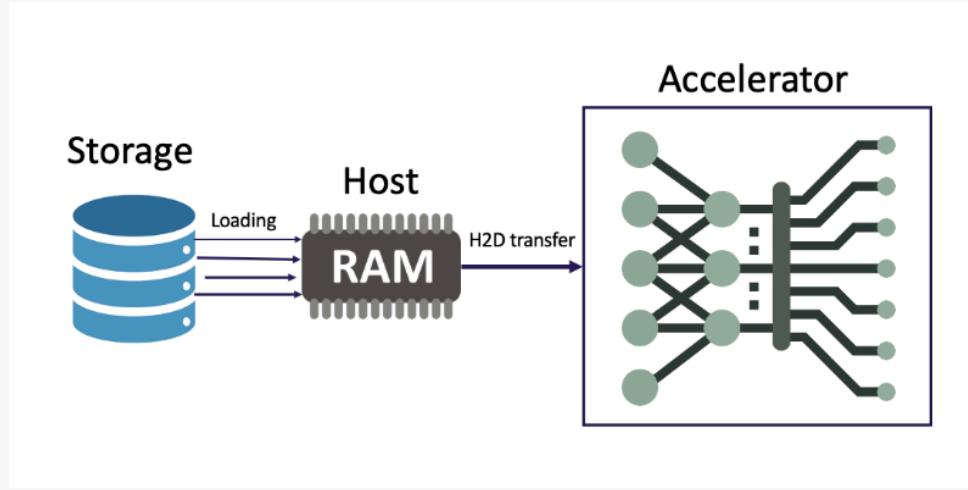
Data Management and I/O for AI

Devarajan, H.; Zheng, H.; Kougkas, A.; Sun, X.-H.; Vishwanath, V. DLIO: A Data-Centric Benchmark for Scientific Deep Learning Applications. In *2021 IEEE/ACM 21st International Symposium on Cluster, Cloud and Internet Computing (CCGrid)*; 2021; pp 81–91.

DLIO Benchmark: https://github.com/argonne-lcf/dlio_benchmark.git

MLPerf Storage: <https://mlcommons.org/en/news/mlperf-storage/>

Deep Learning I/O characteristics

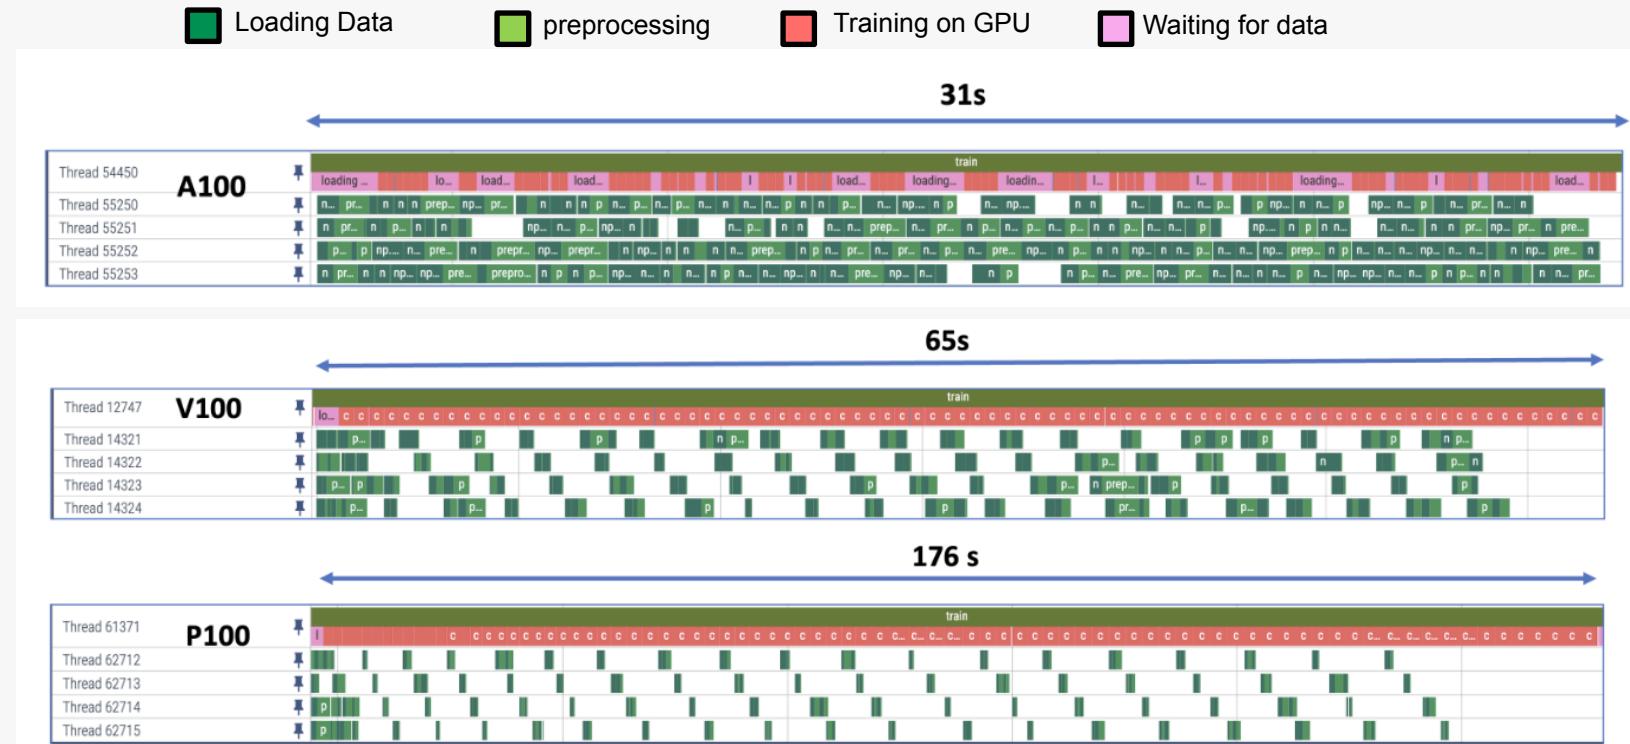


Typical process of AI training. The dataset is loaded from the storage to the host RAM and then feed into the accelerators for training.

Characteristics of I/O for AI applications

- Read intensive
- Metadata intensive
- Small and sparse I/O operations
- Random access
- Complex data format (json, text, key-value store)
- Utilizing storage hierarchy
- Multithreading background I/O

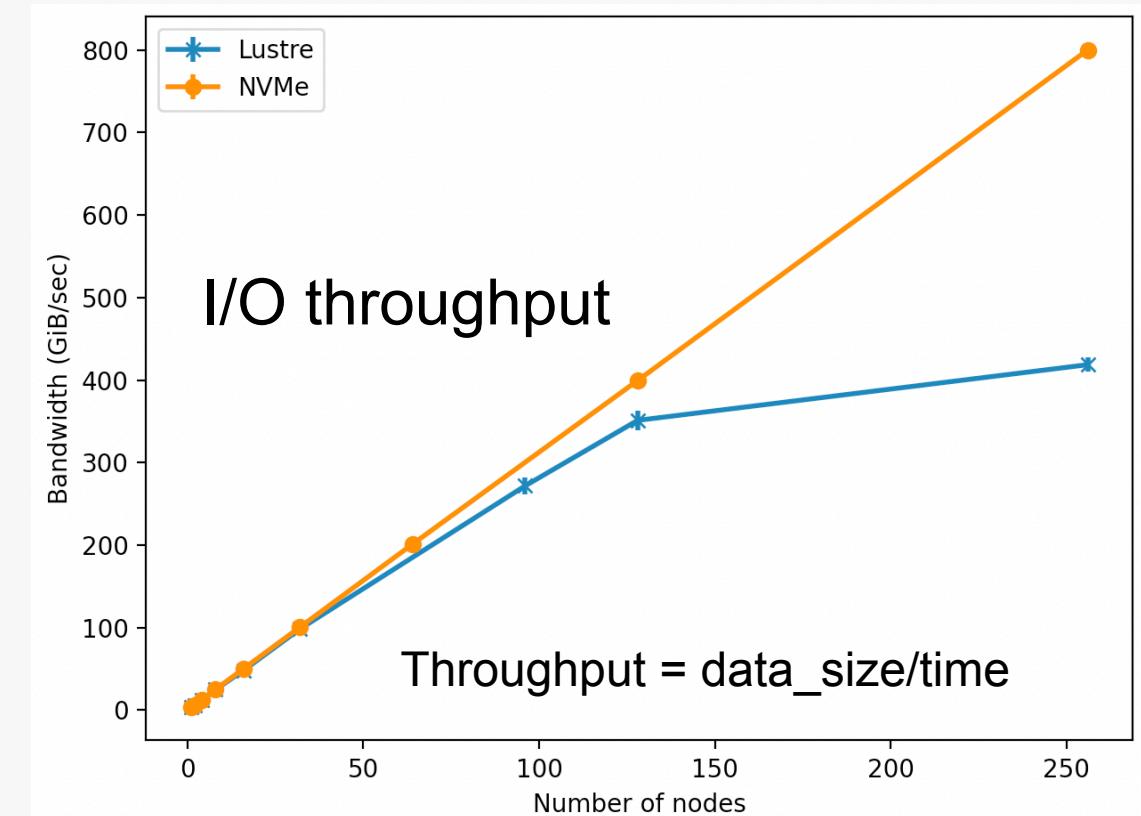
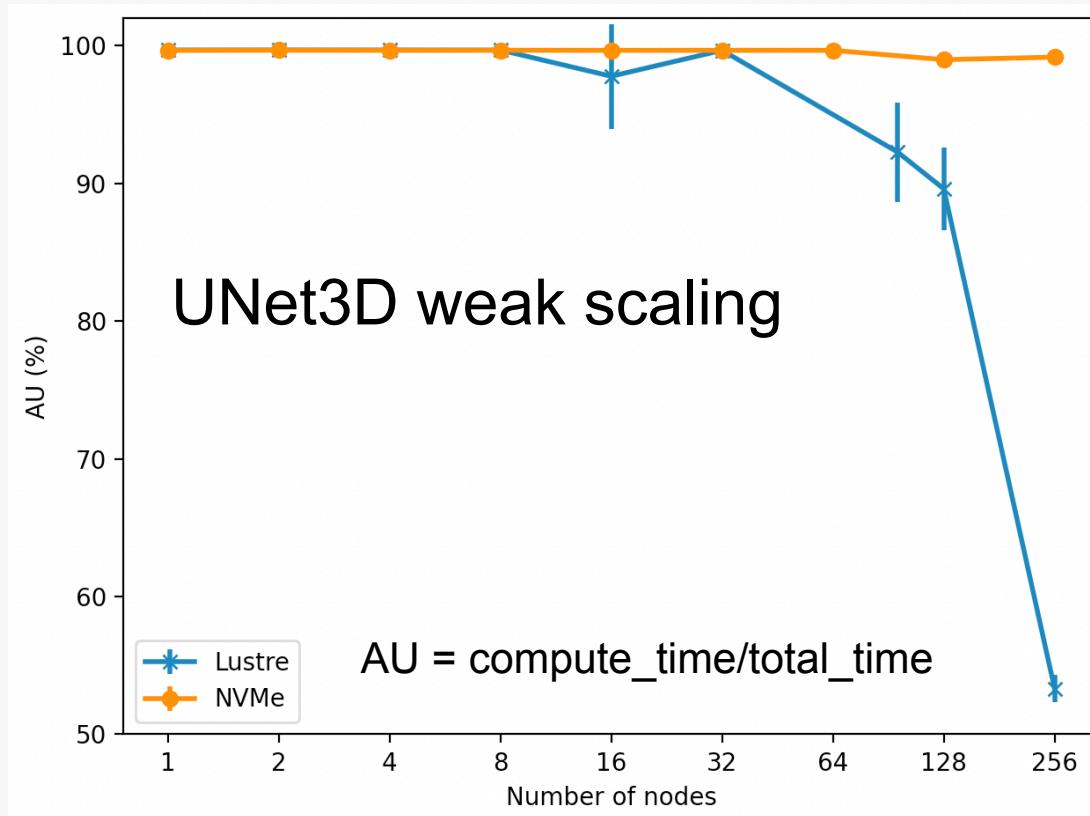
I/O bottleneck for UNet3D workload on fast accelerator



Timeline tracing for training the UNet3D workload on a single GPU on JLSE with GPFS file system. This shows that I/O become a bottleneck for faster accelerator

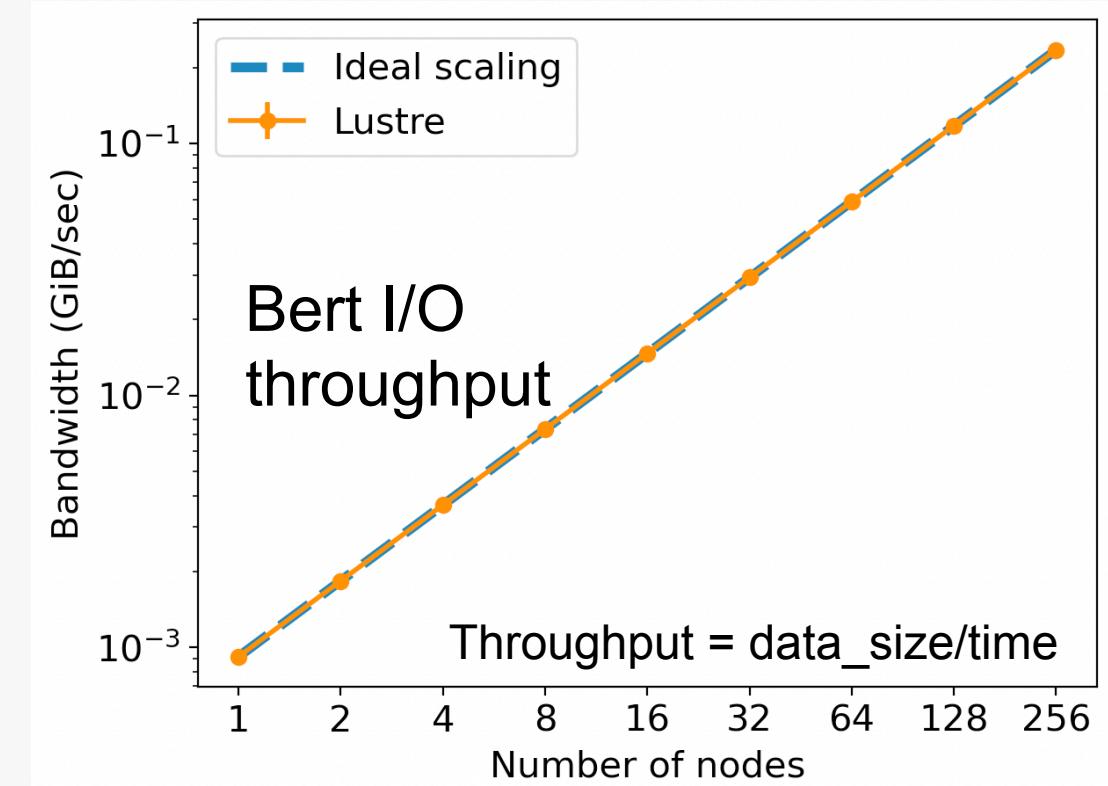
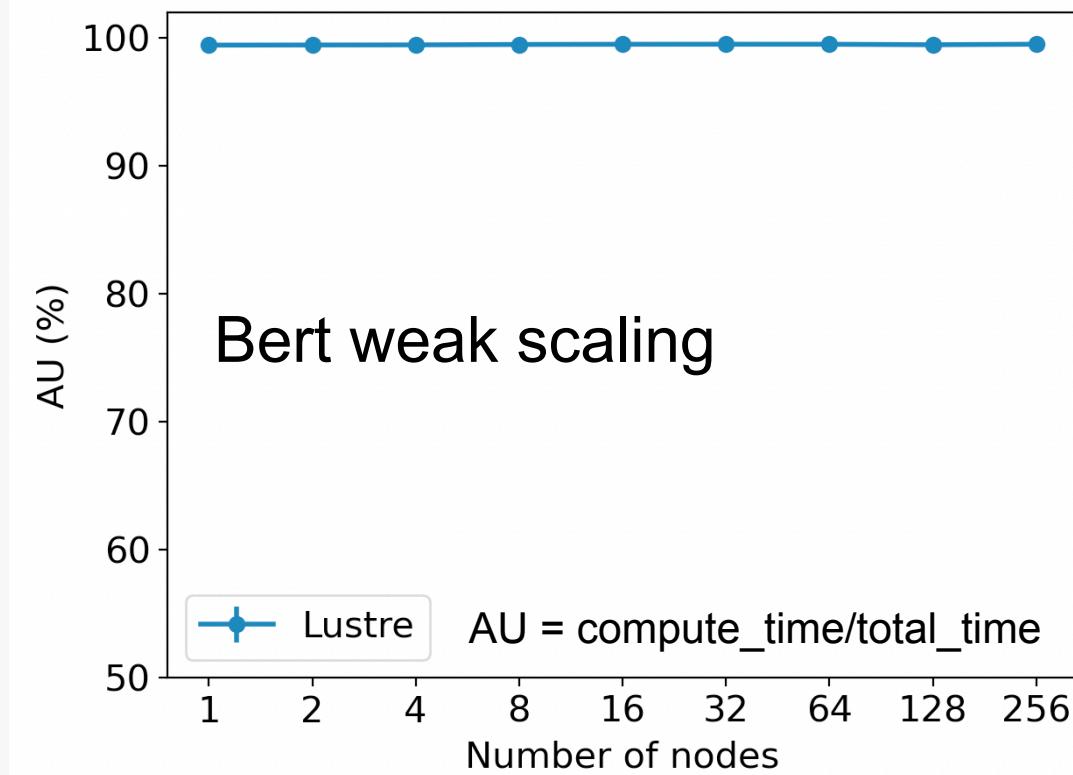
UNet3D Model: https://github.com/mlcommons/training/tree/master/image_segmentation/pytorch

Scaling bottleneck from IO for UNet3D workload (simulated using DLIO Benchmark)



Accelerator utilization (AU) and I/O throughput at different scale on Polaris for UNet3D model, with Lustre file system and NVMe -> staging helps

Ideal scaling for less I/O intensive workload – Bert (simulated using DLIO Benchmark)

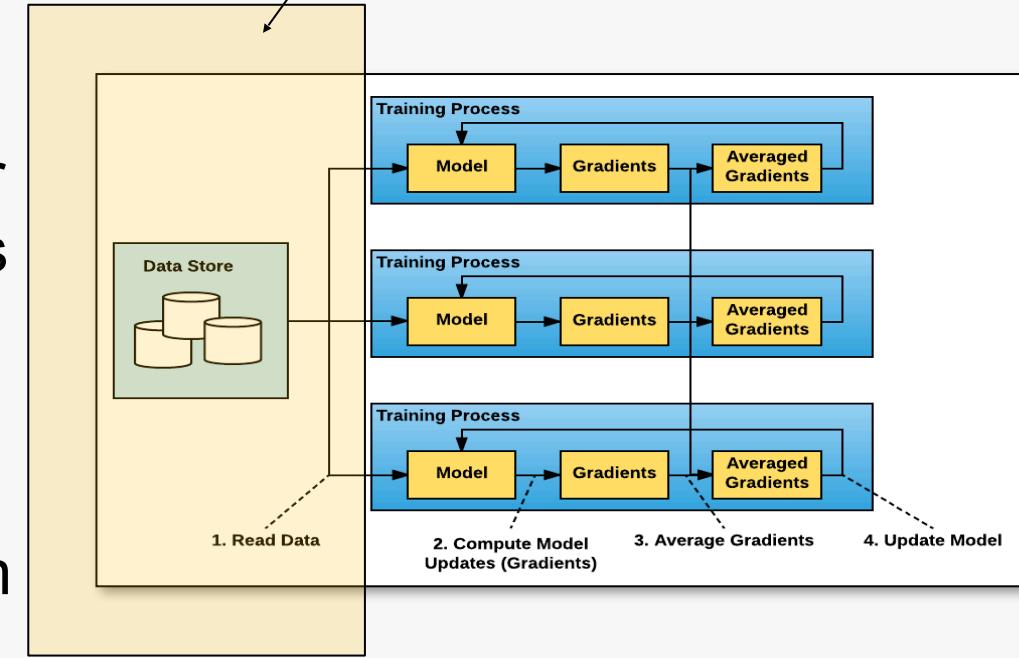


Accelerator utilization (AU) and I/O throughput at different scale on Polaris
for Bert model, with the Lustre file system

Tips for I/O and data management

- Preprocess the raw data (resize, interpolation, etc) into binary format before the training;
- Store the dataset in a reasonable way (file per sample, single shared file, or multiple samples per file)
- Optimal setting (Lustre stripe count, size)
- Remember to shard the dataset;
- Prefetch and caching the data (from disk; from host to device; staging to NVMe, SSDs);
- Use more I/O workers to load data concurrently (e.g., adjust num_workers in TorchDataLoader)

I/O and data management



Streaming I/O using Data Loader

- TensorFlow Data Pipeline
- PyTorch Data Loader
- Nvidia Dali Data Loader

Main take aways

- Distributed training can be done through model parallelism and Data parallelism.
- Data parallelism frameworks: Horovod, DDP (PyTorch only), DeepSpeed (Pytorch only) → similar performance.
- Efficient data management and I/O is crucial for data intensive training

Hands on

```
$ git clone git@github.com:argonne-lcf/ATPESC\_MachineLearning.git
$ cd ATPESC\_MachineLearning/04\_distributedLearning
$ qsub submissions/qsub_polaris.sc
```

The screenshot shows a GitHub repository interface. At the top, the repository path is `argonne-lcf / ATPESC_MachineLearning`. Below the header, there's a navigation bar with links for Code, Issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. The main content area shows the `master` branch with the following commit history:

Name	Last commit message	Last commit date
...		
DDP	commit updating dist dl	1 hour ago
DeepSpeed	commit updating dist dl	1 hour ago
Horovod	fixed learning rate issue	6 minutes ago
figures	rearrange directories for 2023	3 weeks ago
results	rearrange directories for 2023	3 weeks ago
submissions	fixed learning rate issue	6 minutes ago
README.md	fixed learning rate issue	6 minutes ago

Below the commit history, there's a section for the `README.md` file, which contains the following text:

Data Parallel Deep Learning

Author: Huihuo Zheng (huihuo.zheng@anl.gov).

Acknowledgments

- This research used resources of the Argonne Leadership Computing Facility, which is a DOE Office of Science User Facility supported under Contract DE-AC02-06CH11357.
- We gratefully acknowledge the computing resources provided and operated by the Joint Laboratory for System Evaluation (JLSE) at Argonne National Laboratory.



Thank you!

huihuo.zheng@anl.gov